

Reinforcement Learning

Lecture 7: Function approximation

Lecturer: Prof. Dr. Mathias Niepert

Institute for Artificial Intelligence
Machine Learning and Simulation Lab



University of Stuttgart
Germany

imprs-is

June 13, 2024

Outline

1. Introduction
2. Value function approximation
3. Prediction
4. Control
5. Deep Reinforcement Learning


Introduction

Large state spaces

- ▶ Backgammon: 10^{20} states
- ▶ Go: 10^{170} states
- ▶ Continuous spaces:
 - ▶ inverted pendulum
 - ▶ mountain car
 - ▶ helicopter
 - ▶ ...


Small domains: tabular representation

- So far V and Q were just *lookup tables*:



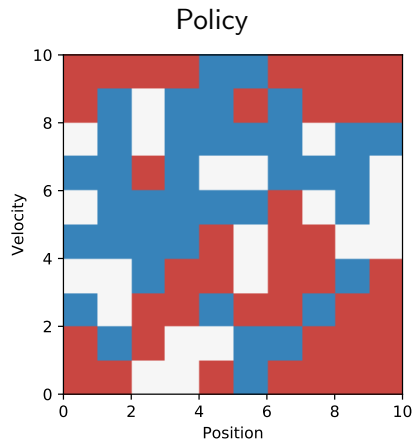

- Problems with large MDPs:
 - too many states to store in memory
 - too slow to learn
- How can we **scale** the tabular solution methods to arbitrarily large state/action spaces?
 - *generalize* from previous encounters of similar states (or state–action pairs)
 - **function approximation** (supervised learning)

Example: mountain car



$a \in \{-1, 0, 1\}$ is the action

State aggregation in mountain car




Value function approximation

Idea of value function approximation

- ▶ Parameterized functional form, with **weights** $\mathbf{w} \in \mathbb{R}^d$:

$$\hat{v}_{\pi}(s, \mathbf{w}) \approx v_{\pi}(s)$$


- ▶ Generally, much less weights than states $d \ll |\mathcal{S}|$
 - ▶ obvious for continuous state spaces
 - ▶ changing single weight, changes value estimate of many states
 - ▶ when one state is updated, change *generalizes* to many states
- ▶ Update \mathbf{w} with MC or TD learning



Function approximators

- ▶ Linear combinations of features
 - ▶ **state aggregation**
 - ▶ **tile coding**
 - ▶ polynomials
 - ▶ radial basis functions (RBFs)
 - ▶ Fourier bases
 - ▶ ...
- ▶ Neural networks
- ▶ Decision trees
- ▶ Non-parametric approaches

Tile coding



Function approximators for RL

- ▶ Differentiable function approximators, e.g.:
 - ▶ Linear combination of features
 - ▶ Neural networks
- ▶ RL specific problems:
 - ▶ non-stationary
 - ▶ non-iid data
 - ▶ bootstrapping
 - ▶ delayed targets

Random variables are *independent and identically distributed* (iid) if they each have the same probability distribution and are mutually independent

Stochastic Gradient Descent (SGD)

- ▶ Approximate value function $\hat{v}(s, \mathbf{w})$
 - ▶ differentiable for all $s \in \mathcal{S}$
- ▶ Weight vector $\mathbf{w} = (w_1, w_2, \dots, w_d)^\top$
 - ▶ \mathbf{w}_t weight vector at time $t = 0, 1, 2, \dots$
- ▶ Gradient of $f(\mathbf{w})$: $\nabla f(\mathbf{w}) = \left(\frac{\partial f(\mathbf{w})}{\partial w_1}, \frac{\partial f(\mathbf{w})}{\partial w_2}, \dots, \frac{\partial f(\mathbf{w})}{\partial w_d} \right)^\top$
- ▶ Do gradient descent by sampling additive parts of the full gradient (i.e., each state consecutively)
- ▶ We can compute our update over smaller sets of inputs

Stochastic Gradient Descent (SGD) (part 2)

- ▶ When we **approximate** the gradient
- ▶ For example

$$\mathcal{L}(\mathbf{w}) = \sum_{n=1}^N \mathcal{L}_n(\mathbf{w})$$

where \mathbf{w} are weights.

- ▶ In machine learning

$$\mathcal{L}(\mathbf{w}) = - \sum_{n=1}^N \log p(y_n \mid \mathbf{x}_n, \mathbf{w})$$

where $\mathbf{x}_n \in \mathbb{R}^D$ are training *inputs*
and y_n are the training *targets*

- ▶ The corresponding update

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \alpha_t \sum_{n=1}^N (\nabla \mathcal{L}_n)(\mathbf{w}_t)$$

- ▶ Often the gradient is too difficult to compute (CPU/GPU expensive)
- ▶ **Mini-batch**: random subset
 - a Large: accurate but costly
 - b Small: noisy but cheap

Stochastic Gradient Descent (SGD) (part 3)

- ▶ *Mean squared Value Error*: $\mathcal{L}(\mathbf{w}) = \sum_{s \in \mathcal{S}} \mu(s) [v_\pi(s) - \hat{v}(s, \mathbf{w})]^2$
- ▶ Adjust \mathbf{w} to reduce the error on sample $S_t \mapsto v_\pi(S_t)$:

$$\begin{aligned} \mathbf{w}_{t+1} &= \mathbf{w}_t - \frac{1}{2} \alpha_t (\nabla \mathcal{L}_t)(\mathbf{w}_t) \\ &= \mathbf{w}_t - \frac{1}{2} \alpha_t \nabla \underbrace{[v_\pi(S_t) - \hat{v}(S_t, \mathbf{w}_t)]^2}_{\text{squared sample error}} \\ &= \mathbf{w}_t + \alpha_t [v_\pi(S_t) - \hat{v}(S_t, \mathbf{w}_t)] \nabla \hat{v}(S_t, \mathbf{w}) \end{aligned}$$

- ▶ α_t is a step size parameter
- ▶ Why not use $\alpha = 1$, thus eliminating full error on sample?

Linear methods

- ▶ Special case where $\hat{v}(\cdot, \mathbf{w})$ is *linear* in the weights
- ▶ **Feature vector** $\mathbf{x}(s)$ represents state s :

$$\mathbf{x}(s) = [x_1(s), \quad x_2(s), \quad \dots, \quad x_d(s)]^\top$$

- ▶ Each component of \mathbf{x} is a feature, examples:
 - ▶ distance of robot to landmarks
 - ▶ piece on a specific location on a chess board
- ▶ Value function is represented as a linear combination of features $\mathbf{x}(s)$:

$$\hat{v}(s, \mathbf{w}) = \mathbf{w}^\top \mathbf{x}(s) = \sum_{i=1}^d w_i x_i(s)$$

- ▶ Gradient is simply $\nabla \hat{v}(s, \mathbf{w}) = \mathbf{x}(s)$

Prediction

Prediction with function approximation

- ▶ We assumed the true value function $v_\pi(S_t)$ is known
- ▶ Substitute target U_t for $v_\pi(s)$

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha[U_t - \hat{v}(S_t, \mathbf{w}_t)]\nabla\hat{v}(S_t, \mathbf{w}_t)$$

- ▶ U_t might be a noisy or bootstrapped approximation of the true value
- ▶ **Monte Carlo:** $U_t = G_t$
- ▶ **TD(0):** $U_t = R_{t+1} + \gamma\hat{v}(S_{t+1}, \mathbf{w}_t)$
- ▶ **TD(λ):** $U_t = G_t^\lambda$

Monte–Carlo with function approximation

- ▶ Target is unbiased by definition:

$$\mathbb{E}[U_t | S_t = s] = \mathbb{E}[G_t | S_t = s] = v_\pi(S_t)$$

- ▶ Training data:

$$\mathcal{D} = \{(S_1, G_1), (S_2, G_2), \dots, (S_{T-1}, G_{T-1}), (S_T, 0)\}$$

- ▶ Using SGD, \mathbf{w} is guaranteed to converge to a *local optimum*
- ▶ MC prediction exhibits local convergence with linear and non-linear function approximation
- ▶ SGD update for sample $S_t \mapsto G_t$:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha[G_t - \hat{v}(S_t, \mathbf{w})]\nabla \hat{v}(S_t, \mathbf{w})$$

Gradient Monte Carlo Algorithm

Gradient Monte Carlo Algorithm for Estimating $\hat{v} \approx v_\pi$

Input: the policy π to be evaluated

Input: a differentiable function $\hat{v} : \mathcal{S} \times \mathbb{R}^d \rightarrow \mathbb{R}$

Algorithm parameter: step size $\alpha > 0$

Initialize value-function weights $\mathbf{w} \in \mathbb{R}^d$ arbitrarily (e.g., $\mathbf{w} = \mathbf{0}$)

Loop forever (for each episode):

 Generate an episode $S_0, A_0, R_1, S_1, A_1, \dots, R_T, S_T$ using π

 Loop for each step of episode, $t = 0, 1, \dots, T - 1$:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [G_t - \hat{v}(S_t, \mathbf{w})] \nabla \hat{v}(S_t, \mathbf{w})$$

TD with function approximation

- ▶ TD-target $U_t = R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w})$ is *biased* sample of the true value $v_\pi(S_t)$
- ▶ Training data:

$$\mathcal{D} = \{(S_1, R_2 + \gamma \hat{v}(S_2, \mathbf{w})), (S_2, R_3 + \gamma \hat{v}(S_3, \mathbf{w})), \dots, (S_{T-1}, R_T)\}$$

- ▶ SGD update for sample $S_t \mapsto G_t$:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}) - \hat{v}(S_t, \mathbf{w})] \nabla \hat{v}(S_t, \mathbf{w})$$

- ▶ Linear TD(0):

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha \left[\underbrace{R_{t+1} + \gamma \mathbf{w}^T \mathbf{x}(S_{t+1})}_{U_t: \text{TD-Target}} - \underbrace{\mathbf{w}^T \mathbf{x}(S_t)}_{\hat{v}: \text{value function}} \right] \mathbf{x}(S_t)$$

Semi-Gradient TD(0) Algorithm

Semi-gradient TD(0) for estimating $\hat{v} \approx v_\pi$

Input: the policy π to be evaluated

Input: a differentiable function $\hat{v} : \mathcal{S}^+ \times \mathbb{R}^d \rightarrow \mathbb{R}$ such that $\hat{v}(\text{terminal}, \cdot) = 0$

Algorithm parameter: step size $\alpha > 0$

Initialize value-function weights $\mathbf{w} \in \mathbb{R}^d$ arbitrarily (e.g., $\mathbf{w} = \mathbf{0}$)

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose $A \sim \pi(\cdot | S)$

 Take action A , observe R, S'


$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R + \gamma \hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w})] \nabla \hat{v}(S, \mathbf{w})$

$S \leftarrow S'$

 until S' is terminal

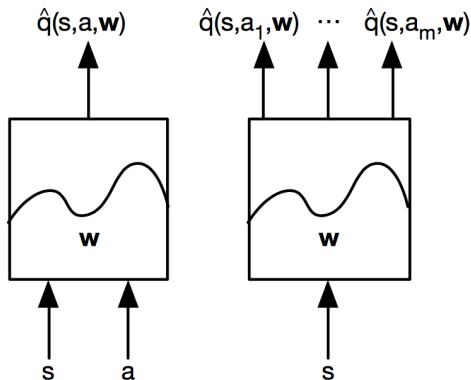
Control

Control with function approximation



- Control via generalized policy iteration (GPI):
 - *policy evaluation*: **approximate** policy evaluation: $\hat{q}(\cdot, \cdot, w) \approx q_\pi$
 - *policy improvement*: ϵ -greedy policy improvement

Types of action–value function approximation



- ▶ Action as input: $\hat{q}(s, a, \mathbf{w}) \approx q_\pi(s, a)$
- ▶ Multiple action–value outputs: $\hat{q}_a(s, \mathbf{w}) \approx q_\pi(s, a)$

Action-value function approximation

- ▶ Approximate action-value function $\hat{q}(s, a, \mathbf{w}) \approx q_\pi(s, a)$
- ▶ Linear case:

$$\hat{q}(s, a, \mathbf{w}) = \mathbf{w}^T \mathbf{x}(s, a) = \sum_{i=1}^d w_i x_i(s, a)$$

$$\nabla \hat{q}(s, a, \mathbf{w}) = \mathbf{x}(s, a)$$

- ▶ Minimize squared error on samples $S_t, A_t \mapsto q_\pi$: $[q_\pi - \hat{q}(S_t, A_t, \mathbf{w})]^2$
- ▶ Use SGD to find local minimum:

$$\begin{aligned} \mathbf{w}_{t+1} &= \mathbf{w}_t - \frac{1}{2} \alpha \nabla [q_\pi(S_t, A_t) - \hat{q}(S_t, A_t, \mathbf{w}_t)]^2 \\ &= \mathbf{w}_t + \alpha [q_\pi(S_t, A_t) - \hat{q}(S_t, A_t, \mathbf{w}_t)] \nabla \hat{q}(S_t, A_t, \mathbf{w}) \end{aligned}$$

Control with function approximation




- ▶ Again, we must substitute target U_t for true action-value $q_\pi(s, a)$:

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha[U_t - \hat{q}(S_t, A_t, \mathbf{w}_t)]\nabla\hat{q}(S_t, A_t, \mathbf{w})$$

- ▶ **Monte Carlo:** $U_t = G_t$
- ▶ **One-step Sarsa:** $U_t = R_{t+1} + \gamma\hat{q}(S_{t+1}, A_{t+1}, \mathbf{w})$

Deep Reinforcement Learning

Recent successes of deep reinforcement learning



Deep (*supervised*) learning

- ▶ **Deep representation** is a composition of *many* functions

$$x \xrightarrow{w_1} h_1 \xrightarrow{w_2} h_2 \xrightarrow{w_3} \dots \xrightarrow{w_n} h_n \xrightarrow{w_{n+1}} y$$

- ▶ Linear transformation and non-linear **activation functions** h_k
- ▶ Weight sharing
 - ▶ **Recurrent neural networks**: across time steps
 - ▶ **Convolutional neural networks**: across spatial (or temporal) regions
- ▶ Weights w optimized by stochastic gradient descent (SGD)
- ▶ Powerful **function approximation** and **representation learning**
 - ▶ finds compact low-dimensional representation (*features*)
- ▶ State-of-the-art for image, text and audio

Naive deep Q-learning

- Q-learning update rule:

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left(r + \gamma \max_a Q(s', a) - Q(s, a) \right)$$

- Q is represented by a neural network with weights \mathbf{w} : $\hat{q}(s, a, \mathbf{w})$
- Loss is the mean-squared TD-error:

$$\mathcal{L}(\mathbf{w}) = \mathbb{E} \left[\left(r + \gamma \max_a \hat{q}(s', a, \mathbf{w}) - \hat{q}(s, a, \mathbf{w}) \right)^2 \right]$$

- Minimize sample errors with SGD

Stability

Naive Q-learning with neural networks oscillates or diverges:

1. Data is non i.i.d!
 - ▶ trajectories
 - ▶ samples are correlated (generated by interaction)
2. Policy changes rapidly with slight changes to Q -values
 - ▶ policy may oscillate
3. Reward range is unknown
 - ▶ gradients can be large
 - ▶ instabilities during back-propagation
4. Maximization bias

Deep Q-networks (DQN)

Deep Q-networks (DQN) address instabilities through:

- ▶ **Experience replay**
 - ▶ store transitions (S_t, A_t, R_t, S_{t+1})
 - ▶ sample random mini-batches
 - ▶ removes correlation, restores i.i.d. property
- ▶ **Target network**
 - ▶ second q network (second set of parameters)
 - ▶ fixed parameters in target network
 - ▶ periodically update target network parameters
- ▶ **Reward clipping/normalization**
 - ▶ clip rewards to $r \in [-1, 1]$
 - ▶ batch normalization

DQN in Atari

“End-to-end” learning:

- ▶ *state*: stack of 4 frames, raw pixels
- ▶ *action*: joystick commands (18 discrete actions)
- ▶ *reward*: change in score

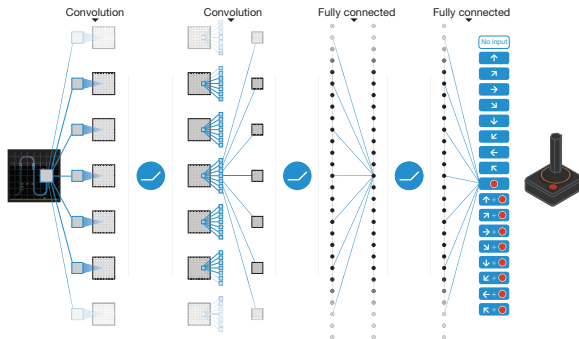


image from *Human-level control through deep reinforcement learning* (Google Deepmind / Nature)

DQN in Atari




image from *Human-level control through deep reinforcement learning* (Google Deepmind / Nature)

Summary

- ▶ When state spaces are **too large** it is not possible to use tables
- ▶ Instead we can use **function approximation** to model v_π and q_π
- ▶ Many function models exist in the literature, the most popular and general are Neural Networks.
- ▶ Stochastic Gradient Descent (SGD) is used to optimize the **Value Function Error**
- ▶ Value function approximation is used for prediction
- ▶ ϵ -greedy or other covering policies can be used for improvement
- ▶ When using a deep-neural network it is important to train in batches using **experience replay** and other techniques to regularize the problem